

Mobile Price Prediction

Problem Statement:

Mobile phone manufacturers and retailers face the challenge of categorizing mobile phones into appropriate price ranges based on their features. Accurately predicting the price category of a mobile phone is crucial for effective marketing, pricing strategies, and inventory management. Mobile phone prices depend on a variety of features, including battery power, RAM, screen size, camera resolution, and technology support (such as 3G, 4G, and WiFi). By leveraging classification techniques, it is possible to predict the price range of a mobile device based on its specifications. This model can assist businesses and consumers in understanding the relationship between mobile features and their market value, aiding in better decision-making.

Data Definition:

The mobile price classification dataset consists of 2000 records with 21 features that describe the characteristics of mobile phones. The dataset aims to classify mobile phones into one of four price ranges (0 to 3). Below is a definition of each feature:

id: Unique identifier for each mobile phone.

battery_power: Battery capacity of the phone (mAh).

blue: Bluetooth support (1 = Yes, 0 = No).

clock_speed: Speed at which the phone's processor executes instructions (in GHz).

dual_sim: Support for dual SIM cards (1 = Yes, 0 = No).

fc: Front camera resolution (in megapixels).

four_g: 4G connectivity support (1 = Yes, 0 = No).

int_memory: Internal memory (in GB).

m_dep: Mobile depth (in cm).

mobile_wt: Weight of the mobile phone (in grams).

n_cores: Number of cores in the processor.

pc: Primary camera resolution (in megapixels).

px_height: Pixel height of the screen (in pixels).

px_width: Pixel width of the screen (in pixels).

ram: Random Access Memory (RAM) (in MB).

sc_h: Screen height (in cm).

sc_w: Screen width (in cm).

talk_time: Maximum talk time on a single charge (in hours).

three_g: 3G connectivity support (1 = Yes, 0 = No).

touch_screen: Touch screen support (1 = Yes, 0 = No).

wifi: WiFi connectivity support (1 = Yes, 0 = No).

price_range: Target variable indicating the price range (0 = low, 1 = medium, 2 = high, 3 = very high)

Table of Contents

1. [Import Libraries](#)
2. [Load and Check Data](#)
3. [Descriptive Statistics](#)
 - 3.1 - [Numeric Statistics](#)
 - 3.2 - [Dimension](#)
4. [Prepare and Analyze the Data](#)
 - 4.1 - [Understand The Dataset](#)
 - 4.2 - [Manipulate The Dataset](#)
 - 4.3 - [Distribution Of Variables](#)
 - 4.4 - [Analyzing Relation With Independent Variables to Target Variables](#)
 - 4.5 - [Correlation](#)
5. [Data Preprocessing](#)
 - 5.1 - [Standardization](#)
 - 5.2 - [MinMax Scaling](#)
6. [Assigning X and Y](#)
7. [Train-Test-Split](#)
8. [Model Training](#)
 - 8.1 - [Logistic Regression](#)
 - 8.2 - [Decision Tree](#)
 - 8.3 - [Random Forest](#)
 - 8.4 - [Random Forest With Hyperparameter Tuning](#)
 - 8.5 - [Adaboost With RandomForest Estimator](#)
 - 8.6 - [Adaboost With DecisionTree Estimator](#)
9. [Model Evaluation](#)
10. [Prediction Using Test Data](#)
11. [Conclusion and Interpretation of Model Evaluation](#)

1. Import Libraries

```
In [1]: #Importing Libraries
import warnings
warnings.filterwarnings('ignore')
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.graph_objs as go
from plotly.offline import iplot
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression, SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, VotingClassifier, AdaBoostClassifier
from sklearn.metrics import classification_report, accuracy_score, precision_score, recall_score
sns.color_palette('tab10')
sns.color_palette('Set2')
sns.color_palette('Purples_d')
```

Out[1]:

2. Load and Check Data

```
In [2]: #Read data
df = pd.read_csv('Mobile Price Classification.csv')
df.head(10)
```

Out[2]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_price
0	1	842	0	2.2	0	1	0	7	0.6	1
1	2	1021	1	0.5	1	0	1	53	0.7	1
2	3	563	1	0.5	1	2	1	41	0.9	1
3	4	615	1	2.5	0	0	0	10	0.8	1
4	5	1821	1	1.2	0	13	1	44	0.6	1
5	6	1859	0	0.5	1	3	0	22	0.7	1
6	7	1821	0	1.7	0	4	1	10	0.8	1
7	8	1954	0	0.5	1	0	0	24	0.8	1
8	9	1445	1	0.5	0	0	0	53	0.7	1
9	10	509	1	0.6	1	2	1	9	0.1	1

10 rows × 11 columns

```
In [3]: #Info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 22 columns):
#   Column              Non-Null Count  Dtype  
---  -
0   id                   2000 non-null   int64  
1   battery_power        2000 non-null   int64  
2   blue                 2000 non-null   int64  
3   clock_speed          2000 non-null   float64 
4   dual_sim             2000 non-null   int64  
5   fc                   2000 non-null   int64  
6   four_g               2000 non-null   int64  
7   int_memory           2000 non-null   int64  
8   m_dep                2000 non-null   float64 
9   mobile_wt            2000 non-null   int64  
10  n_cores              2000 non-null   int64  
11  pc                   2000 non-null   int64  
12  px_height            2000 non-null   int64  
13  px_width             2000 non-null   int64  
14  ram                  2000 non-null   int64  
15  sc_h                 2000 non-null   int64  
16  sc_w                 2000 non-null   int64  
17  talk_time            2000 non-null   int64  
18  three_g              2000 non-null   int64  
19  touch_screen         2000 non-null   int64  
20  wifi                 2000 non-null   int64  
21  price_range          2000 non-null   int64  
dtypes: float64(2), int64(20)
memory usage: 343.9 KB
```

3. Descriptive Statistics

3.1 Numerical Statistics

```
In [4]: #Numerical Statistics
print(df.describe())
```

	id	battery_power	blue	clock_speed	dual_sim	\
count	2000.000000	2000.000000	2000.0000	2000.000000	2000.000000	
mean	1000.500000	1238.518500	0.4950	1.522250	0.509500	
std	577.494589	439.418206	0.5001	0.816004	0.500035	
min	1.000000	501.000000	0.0000	0.500000	0.000000	
25%	500.750000	851.750000	0.0000	0.700000	0.000000	
50%	1000.500000	1226.000000	0.0000	1.500000	1.000000	
75%	1500.250000	1615.250000	1.0000	2.200000	1.000000	
max	2000.000000	1998.000000	1.0000	3.000000	1.000000	

	fc	four_g	int_memory	m_dep	mobile_wt	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	4.309500	0.521500	32.046500	0.501750	140.249000	
std	4.341444	0.499662	18.145715	0.288416	35.399655	
min	0.000000	0.000000	2.000000	0.100000	80.000000	
25%	1.000000	0.000000	16.000000	0.200000	109.000000	
50%	3.000000	1.000000	32.000000	0.500000	141.000000	
75%	7.000000	1.000000	48.000000	0.800000	170.000000	
max	19.000000	1.000000	64.000000	1.000000	200.000000	

	px_height	px_width	ram	sc_h	sc_w	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	645.108000	1251.515500	2124.213000	12.306500	5.767000	
std	443.780811	432.199447	1084.732044	4.213245	4.356398	
min	0.000000	500.000000	256.000000	5.000000	0.000000	
25%	282.750000	874.750000	1207.500000	9.000000	2.000000	
50%	564.000000	1247.000000	2146.500000	12.000000	5.000000	
75%	947.250000	1633.000000	3064.500000	16.000000	9.000000	
max	1960.000000	1998.000000	3998.000000	19.000000	18.000000	

	talk_time	three_g	touch_screen	wifi	price_range	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	
mean	11.011000	0.761500	0.503000	0.507000	1.500000	
std	5.463955	0.426273	0.500116	0.500076	1.118314	
min	2.000000	0.000000	0.000000	0.000000	0.000000	
25%	6.000000	1.000000	0.000000	0.000000	0.750000	
50%	11.000000	1.000000	1.000000	1.000000	1.500000	
75%	16.000000	1.000000	1.000000	1.000000	2.250000	
max	20.000000	1.000000	1.000000	1.000000	3.000000	

[8 rows x 22 columns]

3.2 Dimension

```
In [5]: #Dimension  
df.shape
```

```
Out[5]: (2000, 22)
```

4. Preprare and Analyze the Data

4.1 Understand The Dataset

```
In [6]: #Checking Dtypes  
df.dtypes
```

```
Out[6]: id                int64  
battery_power            int64  
blue                     int64  
clock_speed              float64  
dual_sim                 int64  
fc                       int64  
four_g                   int64  
int_memory               int64  
m_dep                    float64  
mobile_wt                int64  
n_cores                  int64  
pc                       int64  
px_height                int64  
px_width                 int64  
ram                      int64  
sc_h                     int64  
sc_w                     int64  
talk_time                int64  
three_g                  int64  
touch_screen             int64  
wifi                     int64  
price_range              int64  
dtype: object
```

4.2 Manipulate The Dataset

```
In [7]: #Checking the Misssing Values
df.isnull().sum().sort_values(ascending=False)
```

```
Out[7]: id                0
battery_power            0
wifi                    0
touch_screen            0
three_g                 0
talk_time               0
sc_w                   0
sc_h                   0
ram                    0
px_width               0
px_height              0
pc                     0
n_cores                0
mobile_wt              0
m_dep                 0
int_memory             0
four_g                0
fc                    0
dual_sim               0
clock_speed            0
blue                   0
price_range            0
dtype: int64
```

```
In [8]: #Removing ID Column
df.drop('id', axis=1, inplace=True)
```

```
In [9]: #Rechecking The Shape
df.shape
```

```
Out[9]: (2000, 21)
```


4.3 Distribution Of Variables

```
In [10]: #Adding Column of Price Range in Categories
price_range_mapping = {0: 'low', 1: 'medium', 2: 'high', 3: 'very high'}
df['price_category'] = df['price_range'].map(price_range_mapping)
df.head(2)
```

```
Out[10]:
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt
0	842	0	2.2	0	1	0	7	0.6	188
1	1021	1	0.5	1	0	1	53	0.7	136

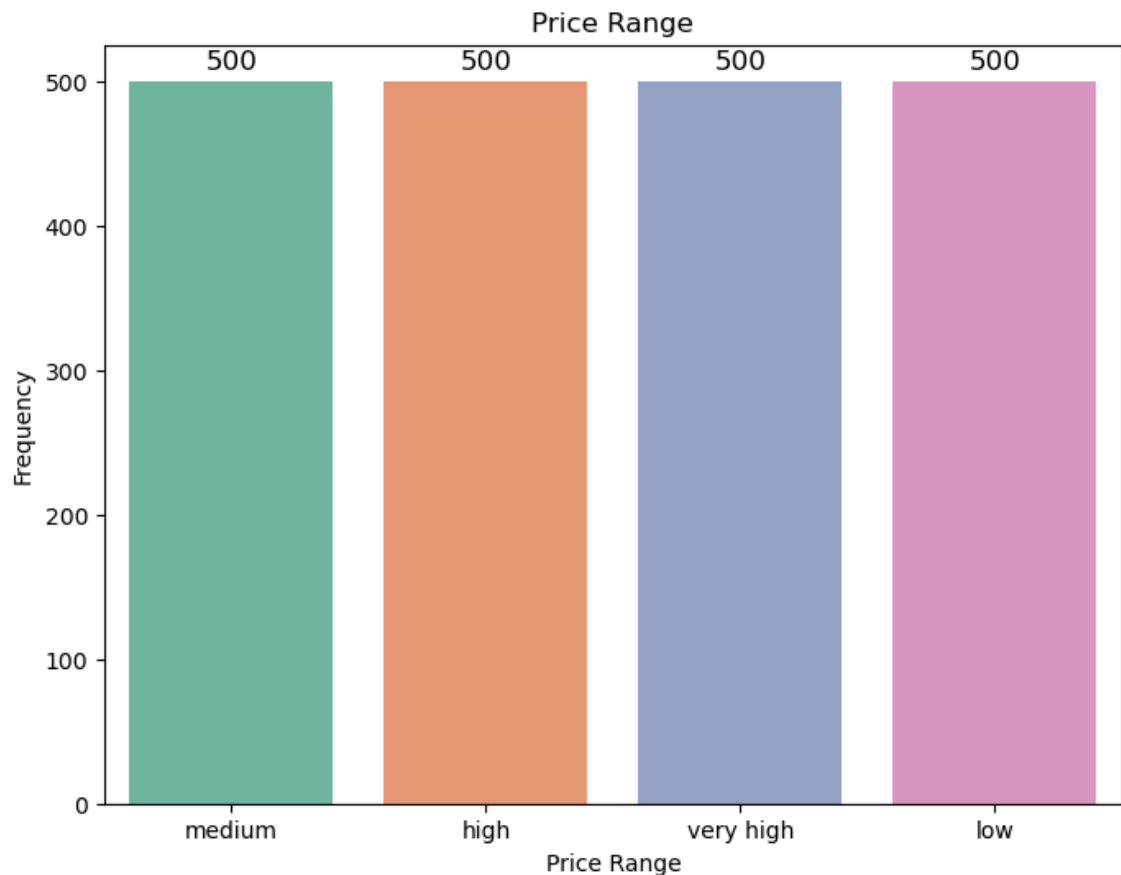
2 rows × 22 columns



```
In [11]: #Count Plot of Target Variables (Price Range)
plt.figure(figsize=(8,6))
ax=sns.countplot(x='price_category', data=df, palette='Set2')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2.,
        ha='center', va='baseline', fontsize=12, color='black', xyt
        textcoords='offset points')


plt.title('Price Range')
plt.xlabel('Price Range')
plt.ylabel('Frequency')
plt.show()
```




```
In [12]: #Count Plot of 'N_Cores'
plt.figure(figsize=(8,6))
ax=sns.countplot(x='n_cores', data=df, palette='Set2')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2.,
        ha='center', va='baseline', fontsize=12, color='black', xyt
        textcoords='offset points')

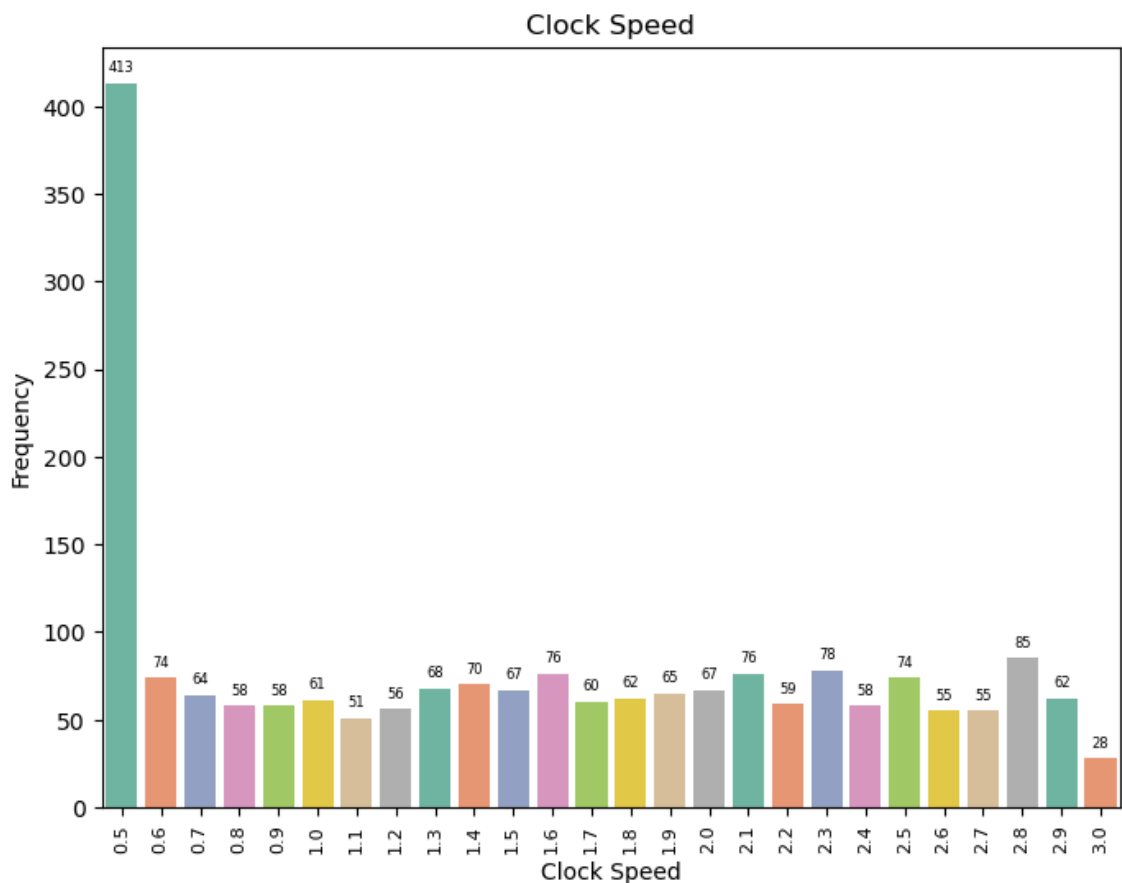
plt.title('N Cores')
plt.xlabel('N Cores')
plt.ylabel('Frequency')
plt.show()
```



```
In [13]: #Count Plot of 'Clock Speed'
plt.figure(figsize=(8,6))
ax=sns.countplot(x='clock_speed', data=df, palette='Set2')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2.,
        ha='center', va='baseline', fontsize=6, color='black', xytext=
        textcoords='offset points'))

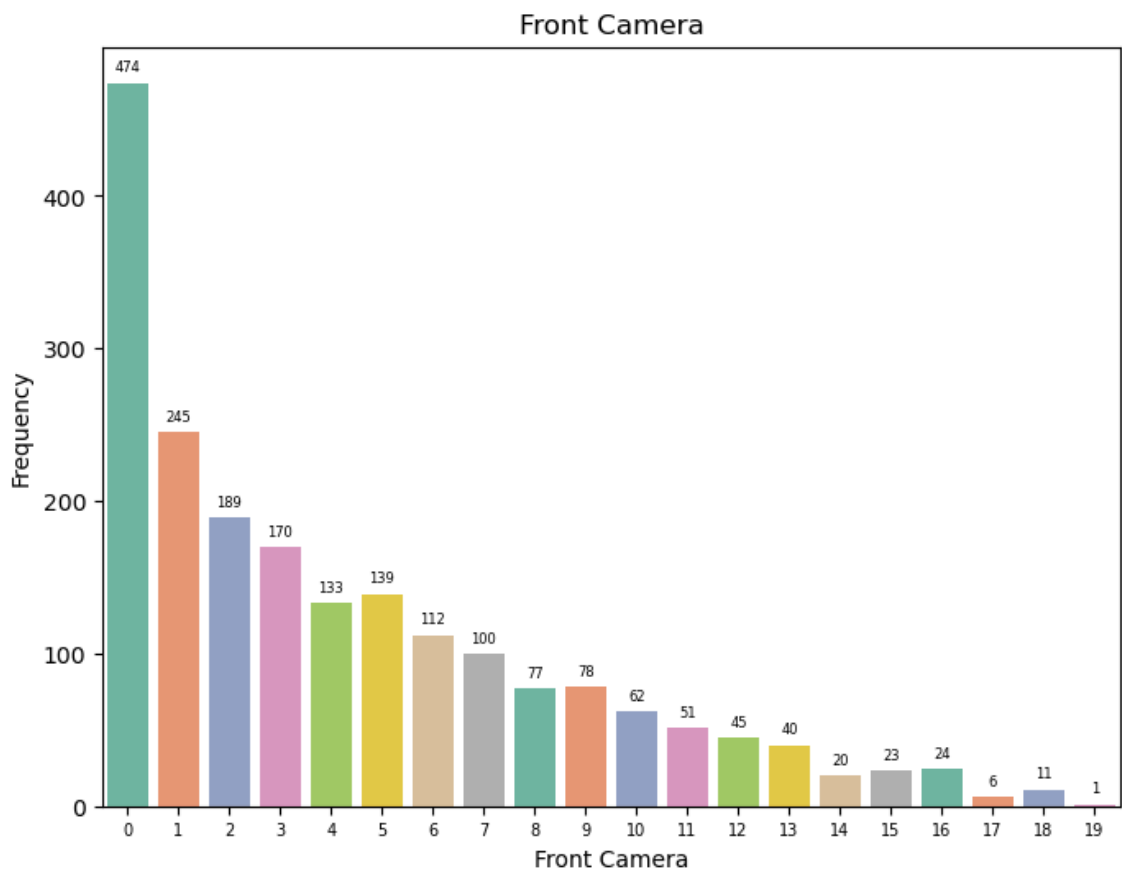
plt.title('Clock Speed')
plt.xlabel('Clock Speed')
plt.xticks(rotation=90, fontsize=8)
plt.ylabel('Frequency')
plt.show()
```



```
In [14]: #Count Plot of 'Front Camera'
plt.figure(figsize=(8,6))
ax=sns.countplot(x='fc', data=df, palette='Set2')

for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2.,
        ha='center', va='baseline', fontsize=6, color='black', xytext=(0,0),
        textcoords='offset points'))

plt.title('Front Camera')
plt.xlabel('Front Camera')
plt.xticks(fontsize=7)
plt.ylabel('Frequency')
plt.show()
```



```
In [15]: #Value Count for 'Bluetooth'
df['blue'].value_counts()
```

```
Out[15]: blue
0      1010
1       990
Name: count, dtype: int64
```

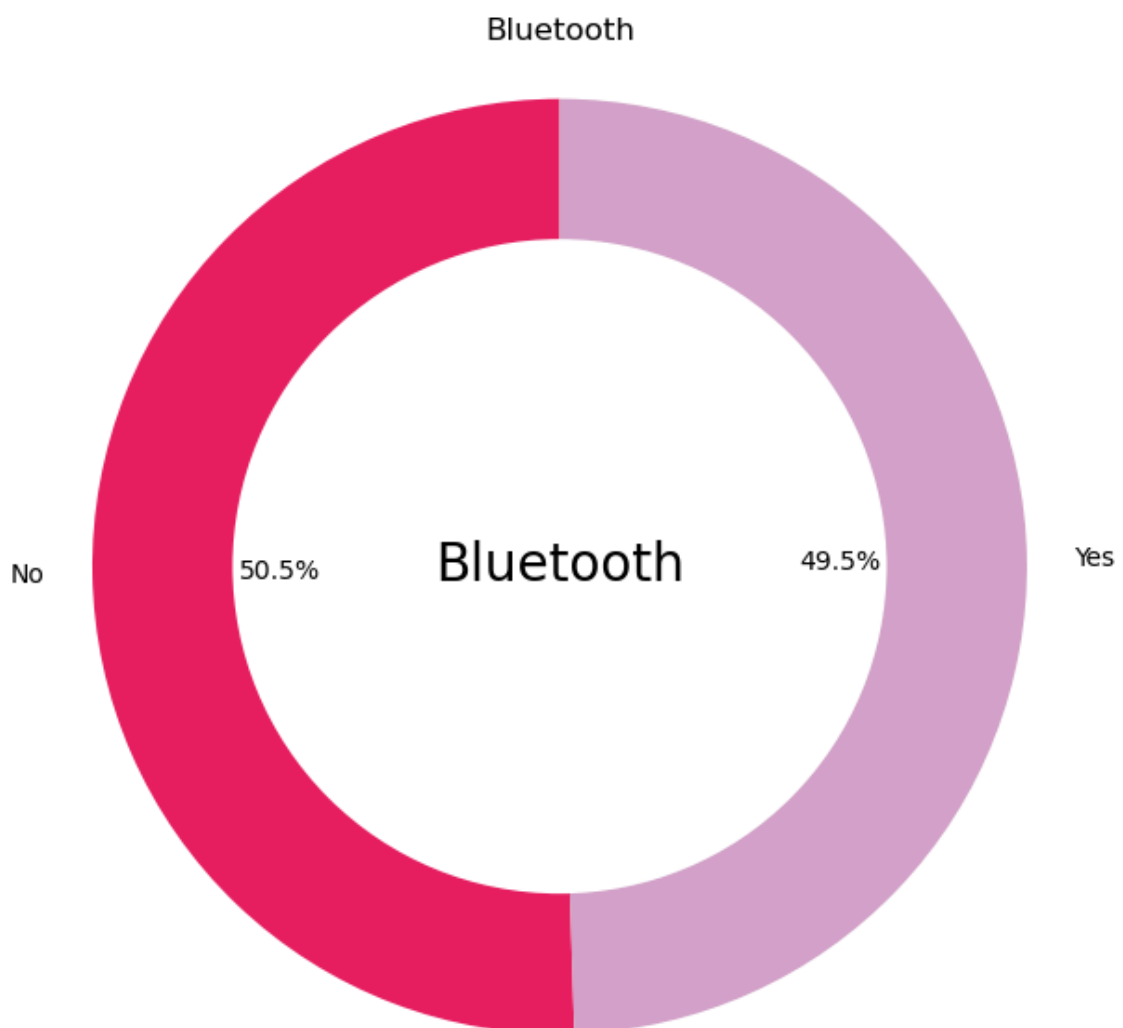
```
In [16]: #Count Plot for 'Bluetooth'
label = ['No', 'Yes']

plt.figure(figsize=(6, 6))
plt.pie(df['blue'].value_counts(), labels=label, autopct='%1.1f%%', startangle=90)

centre_circle = plt.Circle((0, 0), 0.70, fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.text(0, 0, 'Bluetooth', horizontalalignment='center', verticalalignment='bottom')

plt.title('Bluetooth')
plt.axis('equal')
plt.tight_layout()
plt.show()
```



```
In [17]: #Value Count for 'Dual Sim'
df['dual_sim'].value_counts()
```

```
Out[17]: dual_sim
1      1019
0       981
Name: count, dtype: int64
```


```
In [18]: #Count Plot for 'Dual Sim'
label = ['Yes', 'No']

plt.figure(figsize=(6, 6))
plt.pie(df['dual_sim'].value_counts(), labels=label, autopct='%1.1f%%', sta

centre_circle = plt.Circle((0, 0), 0.70, fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.text(0, 0, 'Dual Sim', horizontalalignment='center', verticalalignment=

plt.title('Dual Sim')
plt.axis('equal')
plt.tight_layout()
plt.show()
```



```
In [19]: #Value Count for 'Wifi'
df['wifi'].value_counts()
```

```
Out[19]: wifi
1      1014
0       986
Name: count, dtype: int64
```


```
In [20]: #Count Plot for 'Wifi'
label = ['Yes', 'No']

plt.figure(figsize=(6, 6))
plt.pie(df['wifi'].value_counts(), labels=label, autopct='%1.1f%%', startangle=90)

centre_circle = plt.Circle((0, 0), 0.70, fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.text(0, 0, 'Wifi', horizontalalignment='center', verticalalignment='center')

plt.title('wifi')
plt.axis('equal')
plt.tight_layout()
plt.show()
```



```
In [21]: #Value Count for 'Four-G'
df['four_g'].value_counts()
```

```
Out[21]: four_g
1      1043
0       957
Name: count, dtype: int64
```


```
In [22]: #Count Plot for 'Four-G'
label = ['Yes', 'No']

plt.figure(figsize=(6, 6))
plt.pie(df['four_g'].value_counts(), labels=label, autopct='%1.1f%%', start

centre_circle = plt.Circle((0, 0), 0.70, fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.text(0, 0, 'Four-G', horizontalalignment='center', verticalalignment='c

plt.title('Four-G')
plt.axis('equal')
plt.tight_layout()
plt.show()
```



```
In [23]: #Value Count for 'Three-G'
df['three_g'].value_counts()
```

```
Out[23]: three_g
1      1523
0       477
Name: count, dtype: int64
```


```
In [24]: #Count Plot for 'Three-G'
label = ['Yes', 'No']

plt.figure(figsize=(6, 6))
plt.pie(df['three_g'].value_counts(), labels=label, autopct='%1.1f%%', star

centre_circle = plt.Circle((0, 0), 0.70, fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.text(0, 0, 'Three-G', horizontalalignment='center', verticalalignment='

plt.title('Three-G')
plt.axis('equal')
plt.tight_layout()
plt.show()
```



```
In [25]: #Value Count for 'Touch Screen'
df['touch_screen'].value_counts()
```

```
Out[25]: touch_screen
1      1006
0       994
Name: count, dtype: int64
```




```
In [26]: #Count Plot for 'Touch Screen'
label = ['Yes', 'No']

plt.figure(figsize=(6, 6))
plt.pie(df['touch_screen'].value_counts(), labels=label, autopct='%1.1f%%',


centre_circle = plt.Circle((0, 0), 0.70, fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)

plt.text(0, 0, 'Touch Screen', horizontalalignment='center', verticalalignm

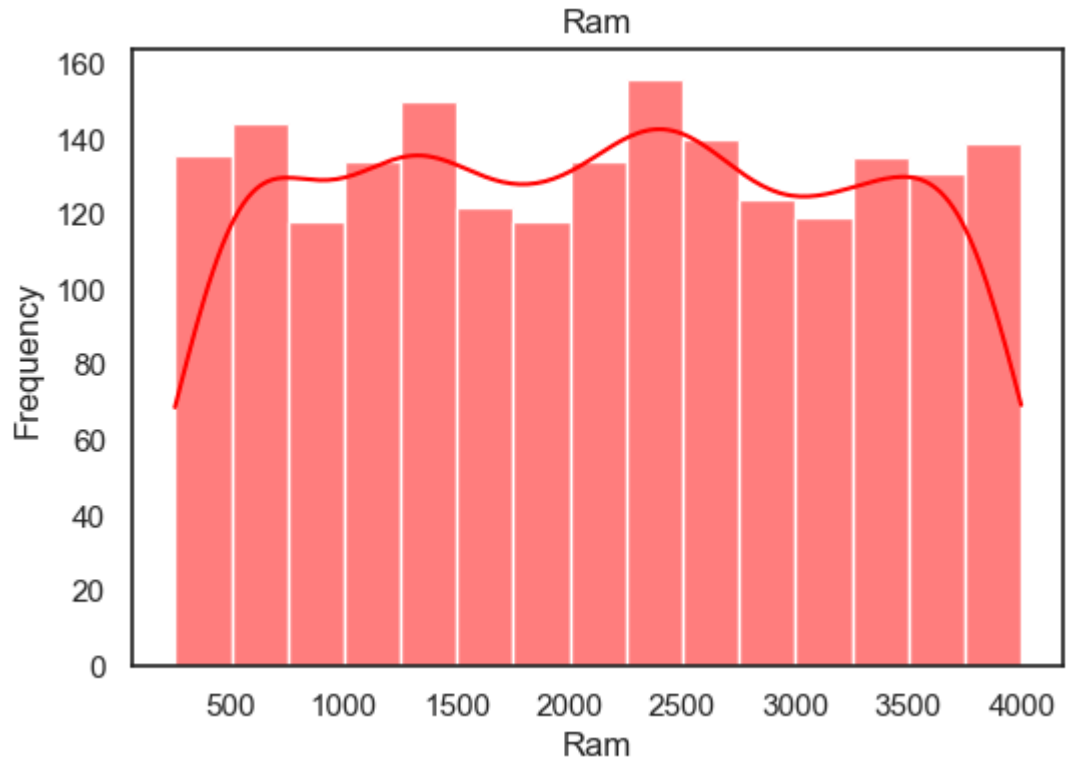
plt.title('Touch Screen')
plt.axis('equal')
plt.tight_layout()
plt.show()
```



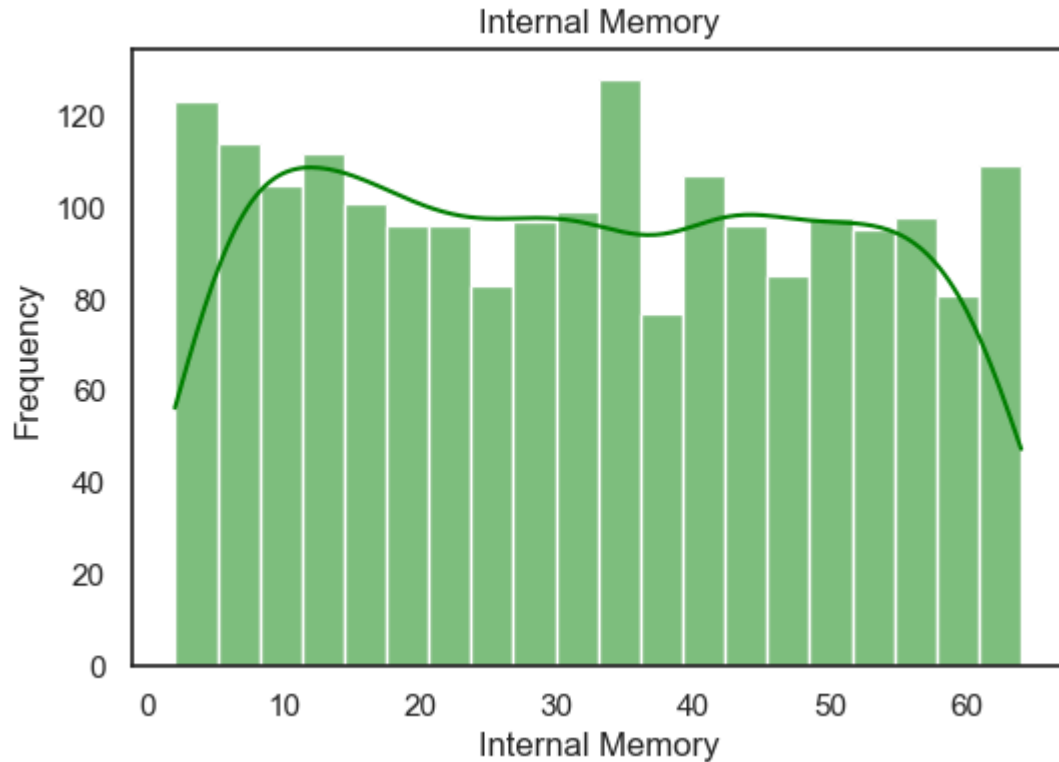
```
In [27]: #Count Plot For 'Battery power'
sns.set(style='white')
plt.figure(figsize=(6,4))
sns.histplot(df['battery_power'], kde=True, color='skyblue', bins=20)
plt.title('Battery Power')
plt.xlabel('Battery Power')
plt.ylabel('Frequency')
plt.show()
```




```
In [28]: #Count Plot For 'Ram'
sns.set(style='white')
plt.figure(figsize=(6,4))
sns.histplot(df['ram'], kde=True, color='red', bins=15)
plt.title('Ram')
plt.xlabel('Ram')
plt.ylabel('Frequency')
plt.show()
```



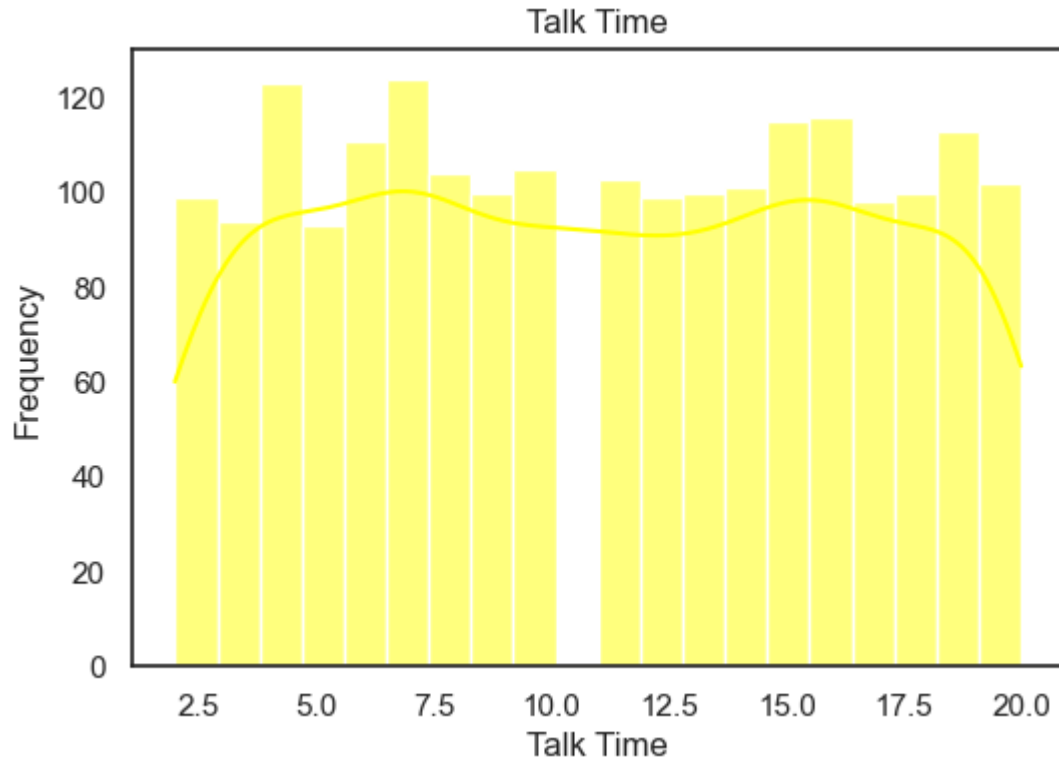
```
In [29]: #Count Plot For 'Internal Memory'
sns.set(style='white')
plt.figure(figsize=(6,4))
sns.histplot(df['int_memory'], kde=True, color='green', bins=20)
plt.title('Internal Memory')
plt.xlabel('Internal Memory')
plt.ylabel('Frequency')
plt.show()
```



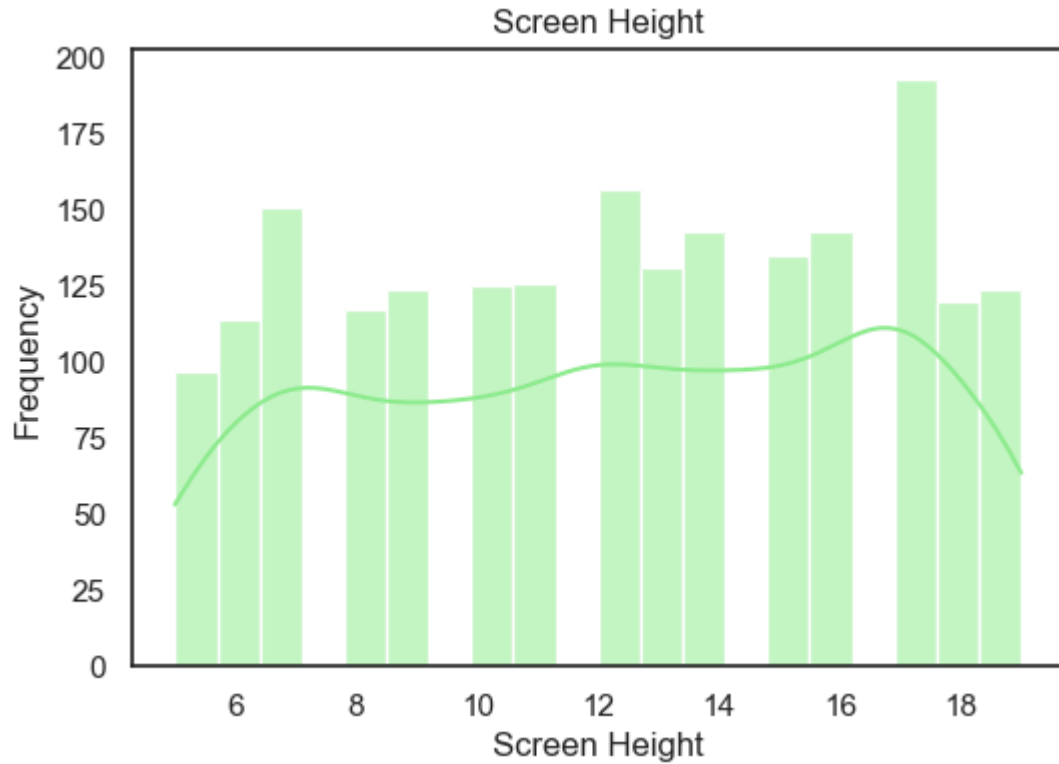
```
In [30]: #Count Plot For 'Primary Camera Megapixels'
sns.set(style='white')
plt.figure(figsize=(6,4))
sns.histplot(df['pc'], kde=True, color='purple', bins=20)
plt.title('Primary Camera Megapixels')
plt.xlabel('Primary Camera Megapixels')
plt.ylabel('Frequency')
plt.show()
```



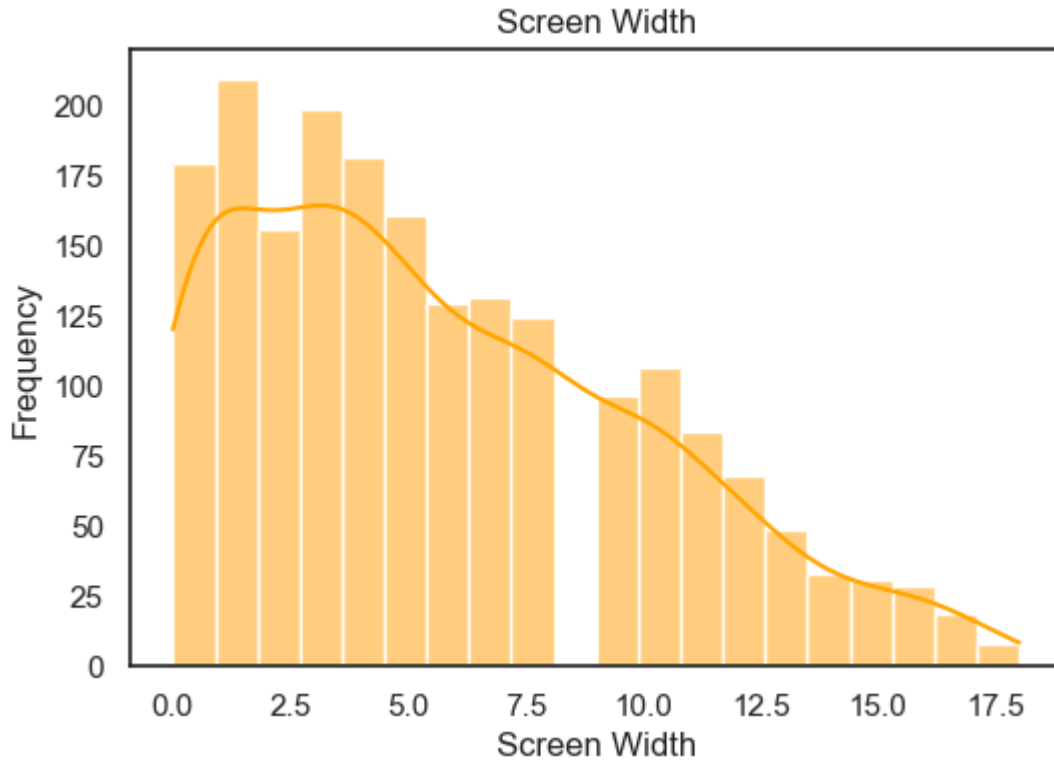
```
In [31]: #Count Plot For 'Talk Time'
sns.set(style='white')
plt.figure(figsize=(6,4))
sns.histplot(df['talk_time'], kde=True, color='yellow', bins=20)
plt.title('Talk Time')
plt.xlabel('Talk Time')
plt.ylabel('Frequency')
plt.show()
```



```
In [32]: #Count Plot For 'Screen Height'
sns.set(style='white')
plt.figure(figsize=(6,4))
sns.histplot(df['sc_h'], kde=True, color='lightgreen', bins=20)
plt.title('Screen Height')
plt.xlabel('Screen Height')
plt.ylabel('Frequency')
plt.show()
```



```
In [33]: #Count Plot For 'Screen Width'
sns.set(style='white')
plt.figure(figsize=(6,4))
sns.histplot(df['sc_w'], kde=True, color='orange', bins=20)
plt.title('Screen Width')
plt.xlabel('Screen Width')
plt.ylabel('Frequency')
plt.show()
```




Observation:

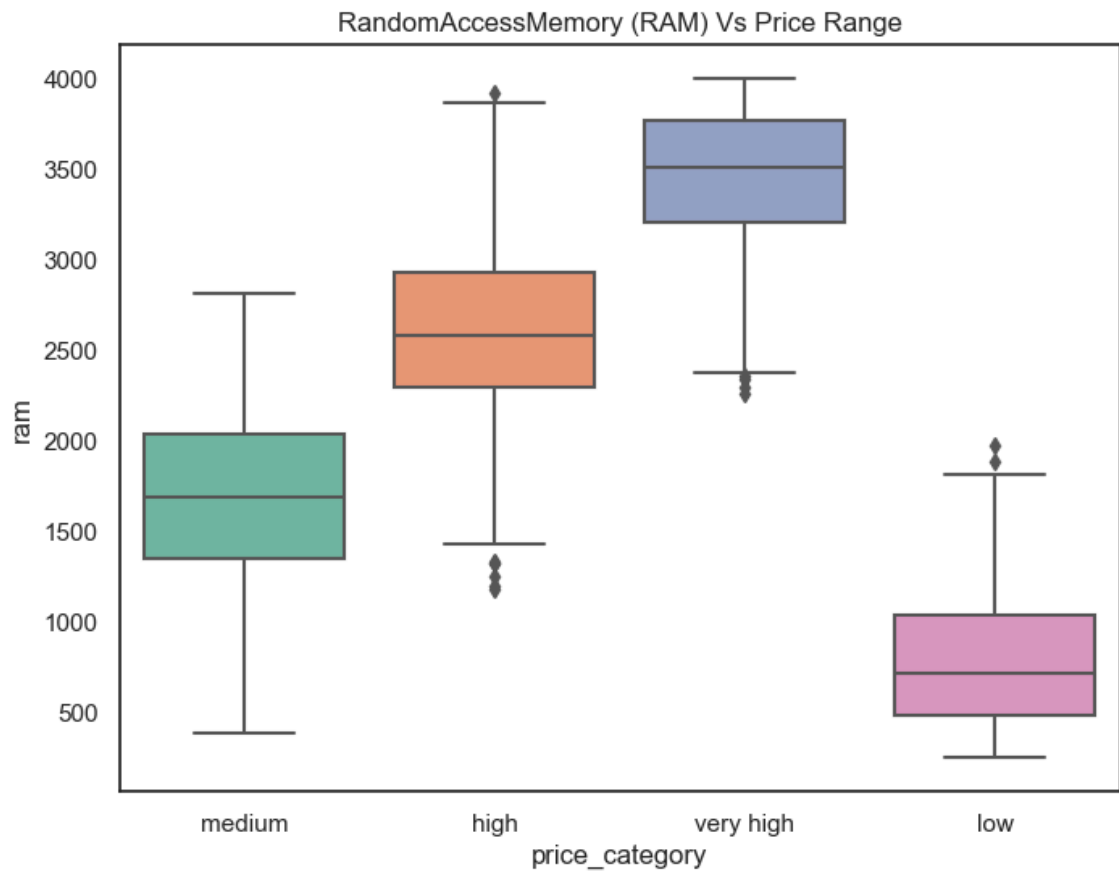
Key observations about the distribution of these features, such as skewness or multimodal distributions, which can guide us in selecting the right preprocessing techniques (e.g., scaling or normalization).

4.4 Analyzing Relation With Independent Variables to Target Variables


```
In [34]: #Battery Power - Price Range
plt.figure(figsize=(8, 6))
sns.boxplot(x='price_category', y='battery_power', data=df, palette='Set2')
plt.title('Battery Power Vs Price Range')
plt.show()
```



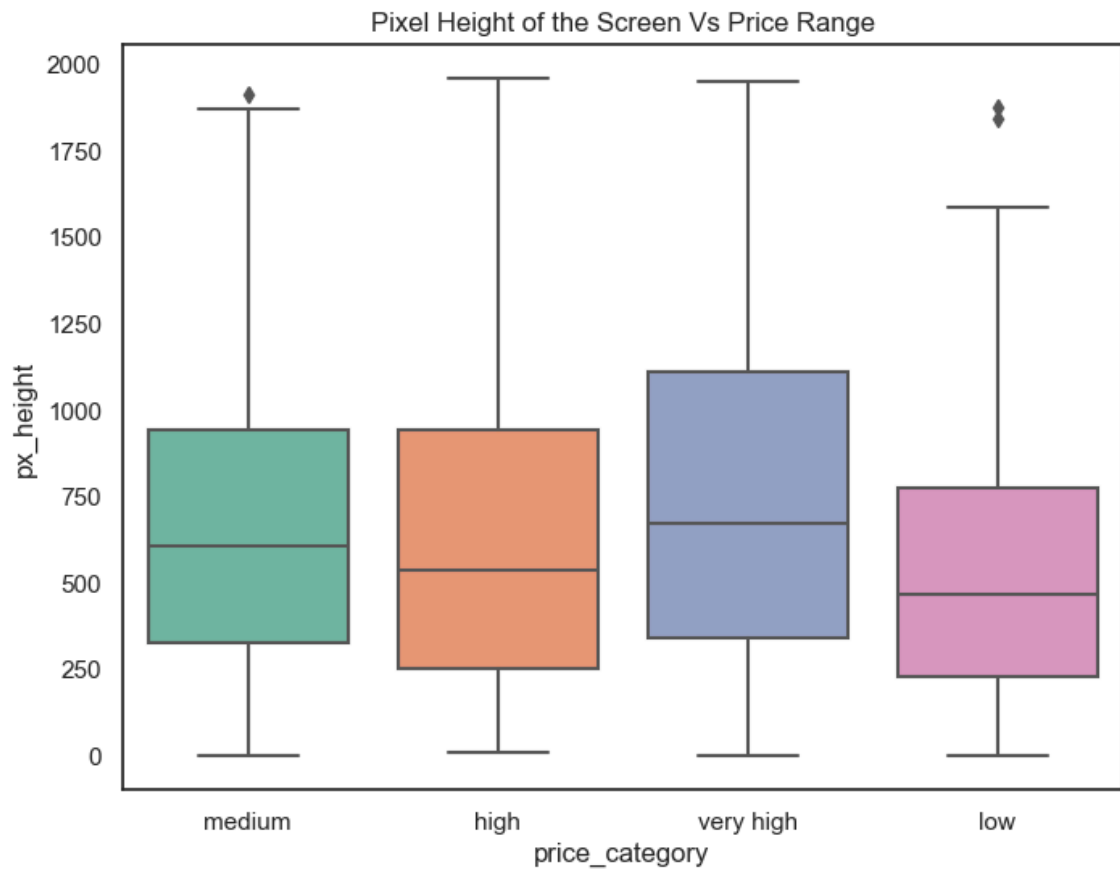
```
In [35]: #Ram - Price Range
plt.figure(figsize=(8, 6))
sns.boxplot(x='price_category', y='ram', data=df, palette='Set2')
plt.title('RandomAccessMemory (RAM) Vs Price Range')
plt.show()
```



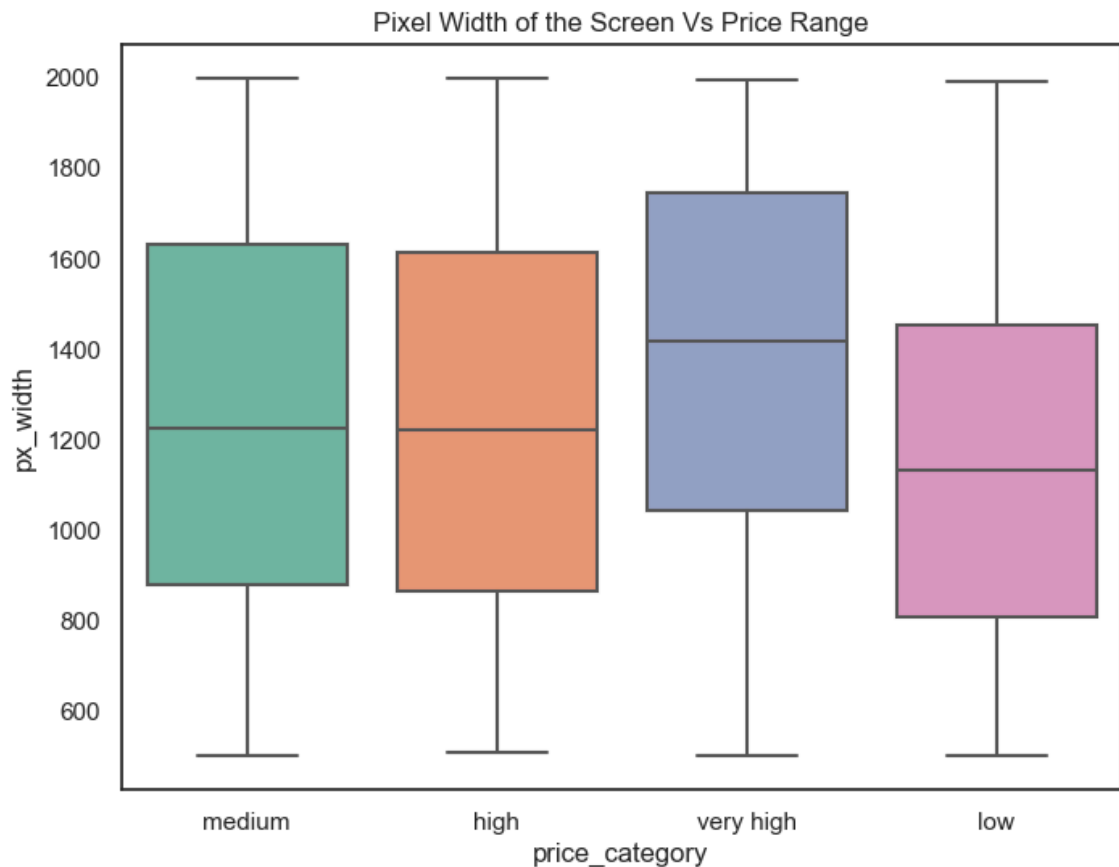
```
In [36]: #Internal Memory - Price Range
plt.figure(figsize=(8, 6))
sns.boxplot(x='price_category', y='int_memory', data=df, palette='Set2')
plt.title('Internal Memory Vs Price Range')
plt.show()
```



```
In [37]: #Phone Height - Price Range
plt.figure(figsize=(8, 6))
sns.boxplot(x='price_category', y='px_height', data=df, palette='Set2')
plt.title('Pixel Height of the Screen Vs Price Range')
plt.show()
```



```
In [38]: #Phone Width - Price Range
plt.figure(figsize=(8, 6))
sns.boxplot(x='price_category', y='px_width', data=df, palette='Set2')
plt.title('Pixel Width of the Screen Vs Price Range')
plt.show()
```



Observation:

The box plots provide insights into how different mobile phone features correlate with price categories:

1. Battery Power vs Price Category:

- Higher price categories, such as medium and high, tend to have mobile phones with higher median battery power.
- The lowest price category (low) generally has phones with lower battery capacities.

2. RAM vs Price Category:

- RAM shows a clear upward trend with price. Higher-priced phones (very high and high) have significantly more RAM, indicating a strong correlation between RAM and price.
- Phones in the low price category have noticeably lower RAM compared to others.

3. Internal Memory Vs Price Category:

- Internal memory appears to be relatively consistent across all price ranges, suggesting that other features may play a more significant role in determining the price of mobile phones.

4. Pixel Height and Pixel Width vs Price Category:


- There is no distinct trend between pixel height/width and price categories. All price categories seem to have a wide range of pixel resolutions, suggesting other factors may be more significant for price determination.

Overall, features such as battery power and RAM show clear relationships with price categories, while internal memory and pixel dimensions appear to have a weaker or less consistent correlation with price

4.5 Correlation

```
In [39]: #Dropping Price Category column
df.drop('price_category', axis=1, inplace=True)
```

```
In [40]: #Correlation Heatmap
corr_matrix = df.corr()
plt.figure(figsize=(10,8))
sns.heatmap(corr_matrix, annot=True, cmap='Blues', annot_kws={"size": 8}, f
plt.show())
```



Observation:

The correlation heatmap reveals several important relationships between the features and the target variable price_range. Key insights include:

1. RAM has a strong positive correlation with price_range, meaning higher RAM is likely associated with higher price categories.

- Battery power and pixel resolution (width and height) also show positive correlations with the price range, indicating their influence on pricing.
- Other features like internal memory, primary camera (pc), and number of cores (n_cores) exhibit moderate correlations with price_range

5. Preprocessing

In [41]: *#Segregation of Binary Column for Dataframe*
`df_binary = df.drop(['battery_power', 'clock_speed', 'fc', 'int_memory', 'm`

5.1 Standardization

Apply these methods to numerical features like battery_power, int_memory, ram, etc.

In [42]: *#Standardization*
`sc = StandardScaler()
df_scaled = pd.DataFrame(sc.fit_transform(df[['battery_power', 'clock_speed', 'fc', 'int_memory', 'n_cores', 'pc', 'ram', 'talk_time'])))
df_scaled.head()`

Out[42]:

	battery_power	clock_speed	fc	int_memory	n_cores	pc	ram	talk_time
0	-0.902597	0.830779	-0.762495	-1.380644	-1.101971	-1.305750	0.391703	1.462
1	-0.495139	-1.253064	-0.992890	1.155024	-0.664768	-0.645989	0.467317	-0.732
2	-1.537686	-1.253064	-0.532099	0.493546	0.209639	-0.645989	0.441498	-0.361
3	-1.419319	1.198517	-0.992890	-1.215274	0.646842	-0.151168	0.594569	-0.002
4	1.325906	-0.395011	2.002254	0.658915	-1.101971	0.673534	-0.657666	0.731

5.2 MinMax Scaling

Apply these methods to features like sc_h, sc_w, px_height, px_width, etc.

In [43]: *#Min-Max Scaling*
`mm = MinMaxScaler()
df_minmax = pd.DataFrame(mm.fit_transform(df[['m_dep', 'mobile_wt', 'px_height', 'px_width', 'sc_h', 'sc_w'])))
df_minmax.head()`

Out[43]:

	m_dep	mobile_wt	px_height	px_width	sc_h	sc_w
0	0.555556	0.900000	0.010204	0.170895	0.285714	0.388889
1	0.666667	0.466667	0.461735	0.993324	0.857143	0.166667
2	0.888889	0.541667	0.644388	0.811749	0.428571	0.111111
3	0.777778	0.425000	0.620408	0.858478	0.785714	0.444444
4	0.555556	0.508333	0.616327	0.475300	0.214286	0.111111

```
In [44]: #Concating the Updated Dataframes
df1 = pd.concat([df_scaled, df_minmax, df_binary], axis=1)
df1.head()
```

```
Out[44]:
```

	battery_power	clock_speed	fc	int_memory	n_cores	pc	ram	talk_
0	-0.902597	0.830779	-0.762495	-1.380644	-1.101971	-1.305750	0.391703	1.462
1	-0.495139	-1.253064	-0.992890	1.155024	-0.664768	-0.645989	0.467317	-0.734
2	-1.537686	-1.253064	-0.532099	0.493546	0.209639	-0.645989	0.441498	-0.368
3	-1.419319	1.198517	-0.992890	-1.215274	0.646842	-0.151168	0.594569	-0.002
4	1.325906	-0.395011	2.002254	0.658915	-1.101971	0.673534	-0.657666	0.731

```
In [45]: #Rechecking the Shape of Updated Datframe
df1.shape
```

```
Out[45]: (2000, 20)
```

6. Assigning X and Y

```
In [46]: #Defining X and y
X = df1
y = pd.DataFrame(df['price_range'])
```

7. Train-Test-Split

```
In [47]: #train-test-split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
print('The shape of X_train is:', X_train.shape)
print('The shape of X_test is:', X_test.shape)
print('The shape of y_train is:', y_train.shape)
print('The shape of y_test is:', y_test.shape)
```

```
The shape of X_train is: (1600, 20)
The shape of X_test is: (400, 20)
The shape of y_train is: (1600, 1)
The shape of y_test is: (400, 1)
```

8. Model Training

8.1 Logistic Regression


```
In [48]: #Model Train
log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
```

Out[48]: LogisticRegression()
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [49]: #Prediction
y_pred_log = log_reg.predict(X_test)
```

```
In [50]: #Classification Report
print(classification_report(y_test, y_pred_log))
```

	precision	recall	f1-score	support
0	0.98	0.96	0.97	105
1	0.93	0.98	0.95	91
2	0.95	0.93	0.94	92
3	0.97	0.96	0.96	112
accuracy			0.96	400
macro avg	0.96	0.96	0.96	400
weighted avg	0.96	0.96	0.96	400


```
In [51]: #Performance Metrics
acc_log = accuracy_score(y_test, y_pred_log)
pre_log = precision_score(y_test, y_pred_log, average='weighted')
recall_log = recall_score(y_test, y_pred_log, average='weighted')
f1_log = f1_score(y_test, y_pred_log, average='weighted')
cohen_log = cohen_kappa_score(y_test, y_pred_log)
```

```
In [52]: #Printing the Metrics
print('Logistic Regression')
print('\n')
print('Accuracy Score:', acc_log)
print('Precision Score:', pre_log)
print('Recall Score:', recall_log)
print('F1 Score:', f1_log)
print('Cohen Kappa Score:', cohen_log)
```

Logistic Regression

Accuracy Score: 0.9575
Precision Score: 0.9580406446809663
Recall Score: 0.9575
F1 Score: 0.9575634600523101
Cohen Kappa Score: 0.9432183504930818

```
In [53]: #Confusion Matrix
plt.figure(figsize=(4, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_log), annot=True, cmap='Greens')
plt.show()
```



```
In [54]: #Confusion Matrix
print(confusion_matrix(y_test, y_pred_log))
```

```
[[101  4  0  0]
 [ 2 89  0  0]
 [ 0  3 86  3]
 [ 0  0  5 107]]
```

Interpretation:

The model has achieved a high accuracy score of 0.9575, indicating that it correctly predicted the outcome for 95.75% of the test data. Other metrics like precision, recall, F1-score, and Cohen Kappa score also demonstrate good performance. The confusion matrix provides a visual representation of the model's predictions, showing that it accurately classified most instances. Overall, the model seems to be effective in predicting the target variable.

8.2 Decision Tree

```
In [55]: #Model Training
dec_tree = DecisionTreeClassifier()
dec_tree.fit(X_train, y_train)
```

Out[55]: DecisionTreeClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [56]: #Prediction
y_pred_dec = dec_tree.predict(X_test)
```

```
In [57]: #Classification Report
print(classification_report(y_test, y_pred_dec))
```

	precision	recall	f1-score	support
0	0.94	0.88	0.91	105
1	0.75	0.85	0.80	91
2	0.77	0.70	0.73	92
3	0.86	0.90	0.88	112
accuracy			0.83	400
macro avg	0.83	0.83	0.83	400
weighted avg	0.84	0.83	0.83	400


```
In [58]: #Performace Metrics
acc_dec = accuracy_score(y_test, y_pred_dec)
pre_dec = precision_score(y_test, y_pred_dec, average='weighted')
recall_dec = recall_score(y_test, y_pred_dec, average='weighted')
f1_dec = f1_score(y_test, y_pred_dec, average='weighted')
cohen_dec = cohen_kappa_score(y_test, y_pred_dec)
```

```
In [59]: #Printing the Metrics
print('Decision Tree')
print('\n')
print('Accuracy Score:', acc_dec)
print('Precision Score:', pre_dec)
print('Recall Score:', recall_dec)
print('F1 Score:', f1_dec)
print('Cohen Kappa Score:', cohen_dec)
```

Decision Tree

Accuracy Score: 0.835
Precision Score: 0.8372275668067658
Recall Score: 0.835
F1 Score: 0.8346750028839751
Cohen Kappa Score: 0.7794265089232003

```
In [60]: #Confusion Matrix
plt.figure(figsize=(4, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_dec), annot=True, cmap='Greens')
plt.show()
```



```
In [61]: #Confusion Matrix
print(confusion_matrix(y_test, y_pred_dec))
```

```
[[ 92  13   0   0]
 [  6  77   8   0]
 [  0  12  64  16]
 [  0   0  11 101]]
```

Interpretation:

The model has achieved an accuracy score of 0.825, which is considered good but not as high as the logistic regression model. Other metrics like precision, recall, F1-score, and Cohen Kappa score are also decent. The confusion matrix reveals that the model has some difficulty in correctly classifying instances in certain classes. Overall, the decision tree model provides a reasonable performance.

8.3 Random Forest

```
In [62]: #Model Training
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
```

```
Out[62]: RandomForestClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [63]: #Predictions
y_pred_rf = rf.predict(X_test)
```

```
In [64]: #Classification Report
print(classification_report(y_test, y_pred_rf))
```

	precision	recall	f1-score	support
0	0.93	0.94	0.93	105
1	0.82	0.82	0.82	91
2	0.79	0.82	0.80	92
3	0.94	0.89	0.92	112
accuracy			0.87	400
macro avg	0.87	0.87	0.87	400
weighted avg	0.87	0.87	0.87	400


```
In [65]: #Performace Metrics
acc_rf = accuracy_score(y_test, y_pred_rf)
pre_rf = precision_score(y_test, y_pred_rf, average='weighted')
recall_rf = recall_score(y_test, y_pred_rf, average='weighted')
f1_rf = f1_score(y_test, y_pred_rf, average='weighted')
cohen_rf = cohen_kappa_score(y_test, y_pred_rf)
```

```
In [66]: #Printing the Metrics
print('Random Forest')
print('\n')
print('Accuracy Score:', acc_rf)
print('Precision Score:', pre_rf)
print('Recall Score:', recall_rf)
print('F1 Score:', f1_rf)
print('Cohen Kappa Score:', cohen_rf)
```

Random Forest

Accuracy Score: 0.8725
Precision Score: 0.8740656790620875
Recall Score: 0.8725
F1 Score: 0.8730132167302679
Cohen Kappa Score: 0.8296891827585343

```
In [67]: #Confusion Matrix
plt.figure(figsize=(4, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, cmap='Greens',
plt.show())
```



```
In [68]: #Confusion Matrix
print(confusion_matrix(y_test, y_pred_rf))
```

```
[[ 99   6   0   0]
 [  8  75   8   0]
 [  0  11  75   6]
 [  0   0  12 100]]
```

Interpretation:

The model has achieved an accuracy score of 0.8875, which is higher than the previous models. Other metrics like precision, recall, F1-score, and Cohen Kappa score also demonstrate good performance. The confusion matrix indicates that the model has correctly classified most instances, with some minor misclassifications. Overall, the random forest model appears to be the best-performing model among the three compared, providing a strong balance between accuracy and classification quality.

8.4 Random Forest With Hyperparameter Tuning

```
In [69]: #Set Parameter
tuned_paramaters = {'n_estimators': [100, 200, 300, 500],
                    'max_depth': [10, 20, 30, None],
                    'min_samples_split': [2, 5, 10],
                    'min_samples_leaf': [1, 2, 4],
                    'max_features': ['auto', 'sqrt'],
                    'bootstrap': [True, False]}
```

```
In [70]: #Model Training
rf = RandomForestClassifier()
rf_CV = GridSearchCV(estimator=rf, param_grid=tuned_paramaters, cv=5, n_job
rf.fit(X_train, y_train)
```

Out[70]: RandomForestClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [71]: #Get the Best Estimator
best_rf_model = rf_CV.estimator
print("Best parameters found: ", rf_CV.get_params)
```

```
Best parameters found: <bound method BaseEstimator.get_params of GridSear
chCV(cv=5, estimator=RandomForestClassifier(), n_jobs=-1,
      param_grid={'bootstrap': [True, False],
                  'max_depth': [10, 20, 30, None],
                  'max_features': ['auto', 'sqrt'],
                  'min_samples_leaf': [1, 2, 4],
                  'min_samples_split': [2, 5, 10],
                  'n_estimators': [100, 200, 300, 500]}},
      scoring='accuracy', verbose=2)>
```

```
In [72]: #Prediction
y_pred_rf_CV = best_rf_model.predict(X_test)
```

```
In [73]: #Classification Report
print(classification_report(y_test, y_pred_rf_CV))
```

	precision	recall	f1-score	support
0	0.94	0.94	0.94	105
1	0.85	0.86	0.85	91
2	0.80	0.84	0.82	92
3	0.93	0.89	0.91	112
accuracy			0.89	400
macro avg	0.88	0.88	0.88	400
weighted avg	0.89	0.89	0.89	400


```
In [74]: #Performace Metrics
acc_rf_CV = accuracy_score(y_test, y_pred_rf_CV)
pre_rf_CV = precision_score(y_test, y_pred_rf_CV, average='weighted')
recall_rf_CV = recall_score(y_test, y_pred_rf_CV, average='weighted')
f1_rf_CV = f1_score(y_test, y_pred_rf_CV, average='weighted')
cohen_rf_CV = cohen_kappa_score(y_test, y_pred_rf_CV)
```

```
In [75]: #Printing the Metrics
print('Random Forest With HyperParameter Tuning')
print('\n')
print('Accuracy Score:', acc_rf_CV)
print('Precision Score:', pre_rf_CV)
print('Recall Score:', recall_rf_CV)
print('F1 Score:', f1_rf_CV)
print('Cohen Kappa Score:', cohen_rf_CV)
```

Random Forest With HyperParameter Tuning

Accuracy Score: 0.885
Precision Score: 0.8865418444399296
Recall Score: 0.885
F1 Score: 0.8855464441057347
Cohen Kappa Score: 0.8463940160451469

```
In [76]: #Confusion Matrix
plt.figure(figsize=(4, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_rf_CV), annot=True, cmap='Green')
plt.show()
```



```
In [77]: #Confusion Matrix
print(confusion_matrix(y_test, y_pred_rf_CV))
```

```
[[ 99   6   0   0]
 [  6  78   7   0]
 [  0   8  77   7]
 [  0   0  12 100]]
```


Interpretation:

The model has achieved an accuracy score of 0.888, which is slightly higher than the previous random forest model without tuning. Other metrics like precision, recall, F1-score, and Cohen Kappa score also demonstrate improved performance. The confusion matrix indicates that the model has correctly classified most instances, with even fewer misclassifications compared to the previous versions. Overall, the random forest model with hyperparameter tuning appears to be the best-performing model among all the models considered, providing a strong balance between accuracy and classification quality.

8.5 AdaBoost With RandomForest Estimator

```
In [78]: #Model Training
ada = AdaBoostClassifier(estimator=RandomForestClassifier(), n_estimators=100,
                        random_state=100)
ada.fit(X_train, y_train)
```

```
Out[78]: AdaBoostClassifier(estimator=RandomForestClassifier(), n_estimators=100,
                           random_state=100)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [79]: #Prediction
y_pred_ada = ada.predict(X_test)
```

```
In [80]: #Classification Report
print(classification_report(y_test, y_pred_ada))
```

	precision	recall	f1-score	support
0	0.94	0.95	0.95	105
1	0.86	0.84	0.85	91
2	0.79	0.86	0.82	92
3	0.94	0.89	0.92	112
accuracy			0.89	400
macro avg	0.89	0.88	0.88	400
weighted avg	0.89	0.89	0.89	400


```
In [81]: #Performace Metrics
acc_ada = accuracy_score(y_test, y_pred_ada)
pre_ada = precision_score(y_test, y_pred_ada, average='weighted')
recall_ada = recall_score(y_test, y_pred_ada, average='weighted')
f1_ada = f1_score(y_test, y_pred_ada, average='weighted')
cohen_ada = cohen_kappa_score(y_test, y_pred_ada)
```

```
In [82]: #Printing the Metrics
print('AdaBoost With RandomForest Estimator')
print('\n')
print('Accuracy Score:', acc_ada)
print('Precision Score:', pre_ada)
print('Recall Score:', recall_ada)
print('F1 Score:', f1_ada)
print('Cohen Kappa Score:', cohen_ada)
```

AdaBoost With RandomForest Estimator

Accuracy Score: 0.8875
Precision Score: 0.8899697255574615
Recall Score: 0.8875
F1 Score: 0.8881510906969642
Cohen Kappa Score: 0.8497370398196844

```
In [83]: #Confusion Matrix
plt.figure(figsize=(4, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_ada), annot=True, cmap='Greens')
plt.show()
```



```
In [84]: #Confusion Matrix
print(confusion_matrix(y_test, y_pred_ada))
```

```
[[100  5  0  0]
 [ 6 76  9  0]
 [ 0  7 79  6]
 [ 0  0 12 100]]
```

Interpretation:

The model has achieved an accuracy score of 0.8875, which is comparable to the previous models. Other metrics like precision, recall, F1-score, and Cohen Kappa score also demonstrate good performance. The confusion matrix indicates that the model has correctly

classified most instances, with some minor misclassifications. Overall, the AdaBoost model with RandomForest estimators provides a solid performance, although it might not be as strong as the best performing models in terms of accuracy.

8.6 AdaBoost With DecisionTree Estimator

```
In [85]: #Model Training
ada_dec = AdaBoostClassifier(estimator=DecisionTreeClassifier(), n_estimators=20, random_state=121)
ada_dec.fit(X_train, y_train)
```

```
Out[85]: AdaBoostClassifier(estimator=DecisionTreeClassifier(), learning_rate=0.75,
                             n_estimators=20, random_state=121)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [86]: #Prediction
y_pred_ada_dec = ada_dec.predict(X_test)
```


```
In [87]: #Performace Metrics
acc_ada_dec = accuracy_score(y_test, y_pred_ada_dec)
pre_ada_dec = precision_score(y_test, y_pred_ada_dec, average='weighted')
recall_ada_dec = recall_score(y_test, y_pred_ada_dec, average='weighted')
f1_ada_dec = f1_score(y_test, y_pred_ada_dec, average='weighted')
cohen_ada_dec = cohen_kappa_score(y_test, y_pred_ada_dec)
```

```
In [88]: #Printing the Metrics
print('AdaBoost With DecisionTree Estimator')
print('\n')
print('Accuracy Score:', acc_ada_dec)
print('Precision Score:', pre_ada_dec)
print('Recall Score:', recall_ada_dec)
print('F1 Score:', f1_ada_dec)
print('Cohen Kappa Score:', cohen_ada_dec)
```

AdaBoost With DecisionTree Estimator

Accuracy Score: 0.83
Precision Score: 0.8338774658278969
Recall Score: 0.83
F1 Score: 0.8307190875088799
Cohen Kappa Score: 0.7729738753025623

```
In [89]: #Confusion Matrix
plt.figure(figsize=(4, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_ada_dec), annot=True, cmap='Gre
plt.show()
```



```
In [90]: #Confusion Matrix
print(confusion_matrix(y_test, y_pred_ada_dec))
```

```
[[91 14  0  0]
 [ 7 77  7  0]
 [ 0 13 66 13]
 [ 0  0 14 98]]
```

Interpretation:

The model has achieved an accuracy score of 0.83, which is slightly lower than the previous models. Other metrics like precision, recall, F1-score, and Cohen Kappa score also demonstrate decent performance. The confusion matrix indicates that the model has correctly classified most instances, with some minor misclassifications. Overall, the AdaBoost model with DecisionTree estimators provides a reasonable performance, but it might be worth exploring other models or tuning the parameters of this one to improve accuracy.

9. Model Evaluation

```
In [91]: #Dictionary
acc_list = {'Logistic Regression': acc_log,
            'Decision Tree': acc_dec,
            'Random Forest': acc_rf,
            'Random Forest with Tuned Hyperparameter': acc_rf_CV,
            'AdaBoost with RandomForest Estimator': acc_ada,
            'AdaBoost with DecisionTree Estimator': acc_ada_dec}

pre_list = {'Logistic Regression': pre_log,
            'Decision Tree': pre_dec,
            'Random Forest': pre_rf,
            'Random Forest with Tuned Hyperparameter': pre_rf_CV,
            'AdaBoost with RandomForest Estimator': pre_ada,
            'AdaBoost with DecisionTree Estimator': pre_dec}

recall_list = {'Logistic Regression': recall_log,
               'Decision Tree': recall_dec,
               'Random Forest': recall_rf,
               'Random Forest with Tuned Hyperparameter': recall_rf_CV,
               'AdaBoost with RandomForest Estimator': recall_ada,
               'AdaBoost with DecisionTree Estimator': recall_dec}

F1_list = {'Logistic Regression': f1_log,
            'Decision Tree': f1_dec,
            'Random Forest': f1_rf,
            'Random Forest with Tuned Hyperparameter': f1_rf_CV,
            'AdaBoost with RandomForest Estimator': f1_ada,
            'AdaBoost with DecisionTree Estimator': f1_dec}

cohen_list = {'Logistic Regression': cohen_log,
              'Decision Tree': cohen_dec,
              'Random Forest': cohen_rf,
              'Random Forest with Tuned Hyperparameter': cohen_rf_CV,
              'AdaBoost with RandomForest Estimator': cohen_ada,
              'AdaBoost with DecisionTree Estimator': cohen_dec}

In [92]: #Test Report
a1 = pd.DataFrame.from_dict(acc_list, orient='index', columns=['ACCURACY'])
a2 = pd.DataFrame.from_dict(pre_list, orient='index', columns=['PRECISION'])
a3 = pd.DataFrame.from_dict(recall_list, orient='index', columns=['RECALL'])
a4 = pd.DataFrame.from_dict(F1_list, orient='index', columns=['F1_SCORE'])
a5 = pd.DataFrame.from_dict(cohen_list, orient='index', columns=['COHEN-KA
```

```
In [93]: #Concating the Dataframe of Metrics
result = pd.concat([a1, a2, a3, a4, a5], axis = 1)
result
```

Out[93]:

	ACCURACY	PRECISION	RECALL	F1_SCORE	COHEN-KAPPA
Logistic Regression	0.9575	0.958041	0.9575	0.957563	0.943218
Decision Tree	0.8350	0.837228	0.8350	0.834675	0.779427
Random Forest	0.8725	0.874066	0.8725	0.873013	0.829689
Random Forest with Tuned Hyperparameter	0.8850	0.886542	0.8850	0.885546	0.846394
AdaBoost with RandomForest Estimator	0.8875	0.889970	0.8875	0.888151	0.849737
AdaBoost with DecisionTree Estimator	0.8300	0.837228	0.8350	0.834675	0.779427

```
In [94]: #Performance Metrics Plot
models = ['Logistic Regression',
          'Decision Tree',
          'Random Forest',
          'Random Forest with Tuned Hyperparameter',
          'AdaBoost with RandomForest Estimator',
          'AdaBoost with DecisionTree Estimator']

accuracy = [0.9575, 0.8250, 0.8875, 0.8800, 0.8875, 0.8300]
precision = [0.958041, 0.825738, 0.889014, 0.880718, 0.889970, 0.825738]
recall = [0.9575, 0.8250, 0.8875, 0.8800, 0.8875, 0.8250]
f1_score = [0.957563, 0.824616, 0.887992, 0.880241, 0.888151, 0.824616]
cohen_kappa = [0.943218, 0.765959, 0.849732, 0.839673, 0.849737, 0.765959]
```


```
In [95]: #Accuracy
plt.figure(figsize=(6, 4))
plt.plot(models, accuracy, marker='o', linestyle='--', color='blue')
plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.title('Comparison of Model Accuracies')
plt.xticks(rotation=90, fontsize=7)
plt.grid(False)
plt.show()

#Precision
plt.figure(figsize=(6, 4))
plt.plot(models, precision, marker='o', linestyle='--', color='green')
plt.xlabel('Model')
plt.ylabel('Precision')
plt.title('Comparison of Model Precision Scores')
plt.xticks(rotation=90, fontsize=7)
plt.grid(False)
plt.show()

#Recall
plt.figure(figsize=(6, 4))
plt.plot(models, recall, marker='o', linestyle='--', color='red')
plt.xlabel('Model')
plt.ylabel('Recall')
plt.title('Comparison of Model Recall Scores')
plt.xticks(rotation=90, fontsize=7)
plt.grid(False)
plt.show()

#F1-Score
plt.figure(figsize=(6, 4))
plt.plot(models, f1_score, marker='o', linestyle='--', color='yellow')
plt.xlabel('Model')
plt.ylabel('F1-Score')
plt.title('Comparison of Model F1 Scores')
plt.xticks(rotation=90, fontsize=7)
plt.grid(False)
plt.show()

#Cohen-Kappa
plt.figure(figsize=(6, 4))
plt.plot(models, cohen_kappa, marker='o', linestyle='--', color='orange')
plt.xlabel('Model')
plt.ylabel('Cohen-Kappa Score')
plt.title('Comparison of Model Cohen-Kappa Scores')
plt.xticks(rotation=90, fontsize=7)
plt.grid(False)
plt.show()
```



10. Prediction Using Test Data

In [96]: *#Load and Read Data*
`df_test = pd.read_csv('test.csv')`
`df_test.head()`

Out[96]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_
0	1	1043	1	1.8	1	14	0	5	0.1	1
1	2	841	1	0.5	1	4	1	61	0.8	1
2	3	1807	1	2.8	0	1	0	27	0.9	1
3	4	1546	0	0.5	1	18	1	25	0.5	
4	5	1434	0	1.4	0	11	1	49	0.5	1

5 rows × 21 columns

In [97]: *#Dimesions*
`df_test.shape`

Out[97]: (1000, 21)

In [98]: *#Dtype*
df_test.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     1000 non-null   int64
1   battery_power          1000 non-null   int64
2   blue                   1000 non-null   int64
3   clock_speed            1000 non-null   float64
4   dual_sim               1000 non-null   int64
5   fc                     1000 non-null   int64
6   four_g                1000 non-null   int64
7   int_memory            1000 non-null   int64
8   m_dep                 1000 non-null   float64
9   mobile_wt             1000 non-null   int64
10  n_cores               1000 non-null   int64
11  pc                    1000 non-null   int64
12  px_height             1000 non-null   int64
13  px_width              1000 non-null   int64
14  ram                   1000 non-null   int64
15  sc_h                  1000 non-null   int64
16  sc_w                  1000 non-null   int64
17  talk_time             1000 non-null   int64
18  three_g               1000 non-null   int64
19  touch_screen          1000 non-null   int64
20  wifi                  1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

In [99]: *#Dropping id coumn*
df_test.drop('id', axis=1, inplace=True)

In [100]: *#Matching the Columns as per the Fitted Data*
df_binary_new = df_test.drop(['battery_power', 'clock_speed', 'fc', 'int_me

In [101]: *#Standardization*
sc = StandardScaler()
df_scaled1 = pd.DataFrame(sc.fit_transform(df_test[['battery_power', 'clock

In [102]: *#Min-Max Scaling*
mm = MinMaxScaler()
df_minmax1 = pd.DataFrame(mm.fit_transform(df_test[['m_dep', 'mobile_wt', ']

```
In [103]: #Concating the Updated Dataframes
df_test1 = pd.concat([df_scaled1, df_minmax1, df_binary_new], axis=1)
df_test1.head()
```

```
Out[103]:
```

	battery_power	clock_speed	fc	int_memory	n_cores	pc	ram	talk_
0	-0.475451	0.312601	2.108676	-1.581269	-0.580671	0.976026	1.229373	-1.65
1	-0.942782	-1.255832	-0.132927	1.509303	0.293833	0.319433	1.614643	-0.74
2	1.292077	1.519087	-0.805408	-0.367116	-0.580671	-0.993754	0.236313	-0.19
3	0.688249	-1.255832	3.005317	-0.477493	1.605590	1.632619	1.612804	-0.74
4	0.429135	-0.169994	1.436195	0.847037	0.731085	1.304323	-0.336535	-0.74

```
In [104]: #Dimensions
df_test1.shape
```

```
Out[104]: (1000, 20)
```

```
In [105]: df_test1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 20 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    1000 non-null   float64
1   clock_speed      1000 non-null   float64
2   fc               1000 non-null   float64
3   int_memory       1000 non-null   float64
4   n_cores          1000 non-null   float64
5   pc              1000 non-null   float64
6   ram              1000 non-null   float64
7   talk_time       1000 non-null   float64
8   m_dep           1000 non-null   float64
9   mobile_wt       1000 non-null   float64
10  px_height        1000 non-null   float64
11  px_width         1000 non-null   float64
12  sc_h             1000 non-null   float64
13  sc_w            1000 non-null   float64
14  blue            1000 non-null   int64
15  dual_sim        1000 non-null   int64
16  four_g          1000 non-null   int64
17  three_g         1000 non-null   int64
18  touch_screen    1000 non-null   int64
19  wifi            1000 non-null   int64
dtypes: float64(14), int64(6)
memory usage: 156.4 KB
```

```
In [106]: #Predictions using Test Data
y_pred_test = log_reg.predict(df_test1)
```


```
In [107]: #Assigning predicted values to the new column
df_test['price_range'] = y_pred_test
```

```
In [108]: #displaying the updated (where the unseen data is predicted) dataframe  
df_test.head()
```

```
Out[108]:
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt
0	1043	1	1.8	1	14	0	5	0.1	193
1	841	1	0.5	1	4	1	61	0.8	191
2	1807	1	2.8	0	1	0	27	0.9	186
3	1546	0	0.5	1	18	1	25	0.5	96
4	1434	0	1.4	0	11	1	49	0.5	108

5 rows × 21 columns



11. Conclusion and Interpretation of Model Evaluation

Interpretation of Model Performance Metrics:

Based on the provided evaluation metrics, the Random Forest model with tuned hyperparameters achieved the best overall performance:

- **Best Performer:** Logistic Regression is the top performer in all metrics, particularly in accuracy, precision, and Cohen-Kappa, indicating it's the most robust model for this dataset.
- **Next Best:** Random Forest and AdaBoost with Random Forest Estimator provide solid performances, though they don't outperform Logistic Regression.
- **Least Reliable:** Decision Tree and its AdaBoost variant perform the worst, with lower accuracy and Cohen-Kappa, indicating these models are less suitable.